

Math Functions

Torch provides MATLAB-like functions for manipulating `Tensor` objects. Functions fall into several types of categories:

- Constructors like `zeros` , `ones` ;
- Extractors like `diag` and `triu` ;
- Element-wise mathematical operations like `abs` and `pow` ;
- BLAS operations;
- Column or row-wise operations like `sum` and `max` ;
- Matrix-wide operations like `trace` and `norm` ;
- Convolution and cross-correlation operations like `conv2` ;
- Basic linear algebra operations like `eig` ;
- Logical operations on `Tensor` s.

By default, all operations allocate a new `Tensor` to return the result.

However, all functions also support passing the target `Tensor` (s) as the first argument(s), in which case the target `Tensor` (s) will be resized accordingly and filled with result.

This property is especially useful when one wants have tight control over when memory is allocated.

The *Torch* package adopts the same concept, so that calling a function directly on the `Tensor` itself using an object-oriented syntax is equivalent to passing the `Tensor` as the optional resulting `Tensor` .

The following two calls are equivalent.

```
torch.log(x, x)
x.log()
```

Similarly, `torch.conv2` function can be used in the following manner.

```
> x = torch.rand(100, 100)
> k = torch.rand(10, 10)
> res1 = torch.conv2(x, k)    -- case 1

> res2 = torch.Tensor()
> torch.conv2(res2, x, k)    -- case 2

> res2:dist(res1)
0
```

The advantage of second case is, same `res2` `Tensor` can be used successively in a loop without any new allocation.

```
-- no new memory allocations...
> for i = 1, 100 do
    torch.conv2(res2, x, k)
end

> res2:dist(res1)
0
```

Construction or extraction functions

`[res] torch.cat([res,] x_1, x_2, [dimension])`

`[res] torch.cat([res,] {x_1, x_2, ...}, [dimension])`

`x = torch.cat(x_1, x_2, [dimension])` returns a `Tensor` `x` which is the concatenation of `Tensor` s `x_1` and `x_2` along dimension `dimension`.

If `dimension` is not specified or if it is `-1`, it is the maximum last dimension over all input tensors, except if all tensors are empty, then it is `1`.

The other dimensions of `x_1` and `x_2` have to be equal.

Also supports arrays with arbitrary numbers of `Tensor` s as inputs.

Empty tensors are ignored during catting, and thus do not throw an error. Performing cat on empty tensors only will always result in an empty tensor.

Examples:

```
> torch.cat(torch.ones(3), torch.zeros(2))
1
1
```

```

1
0
0
[torch.DoubleTensor of size 5]

> torch.cat(torch.ones(3, 2), torch.zeros(2, 2), 1)
1  1
1  1
1  1
0  0
0  0
[torch.DoubleTensor of size 5x2]

> torch.cat(torch.ones(2, 2), torch.zeros(2, 2), 1)
1  1
1  1
0  0
0  0
[torch.DoubleTensor of size 4x2]

> torch.cat(torch.ones(2, 2), torch.zeros(2, 2), 2)
1  1  0  0
1  1  0  0
[torch.DoubleTensor of size 2x4]

> torch.cat(torch.cat(torch.ones(2, 2), torch.zeros(2, 2), 1),
torch.rand(3, 2), 1)
1.0000  1.0000
1.0000  1.0000
0.0000  0.0000
0.0000  0.0000
0.3227  0.0493
0.9161  0.1086
0.2206  0.7449
[torch.DoubleTensor of size 7x2]

> torch.cat({torch.ones(2, 2), torch.zeros(2, 2), torch.rand(3,
2)}, 1)
1.0000  1.0000
1.0000  1.0000
0.0000  0.0000
0.0000  0.0000
0.3227  0.0493
0.9161  0.1086
0.2206  0.7449

```

```
[torch.DoubleTensor of size 7x2]

> torch.cat({torch.Tensor(), torch.rand(3, 2)}, 1)
0.3227  0.0493
0.9161  0.1086
0.2206  0.7449
[torch.DoubleTensor of size 3x2]
```

[res] torch.diag([res,] x [,k])

`y = torch.diag(x)` when `x` is of dimension 1 returns a diagonal matrix with diagonal elements constructed from `x`.

`y = torch.diag(x)` when `x` is of dimension 2 returns a `Tensor` of dimension 1 with elements constructed from the diagonal of `x`.

`y = torch.diag(x, k)` returns the `k`-th diagonal of `x`, where `k = 0` is the main diagonal, `k > 0` is above the main diagonal and `k < 0` is below the main diagonal.

[res] torch.eye([res,] n [,m])

`y = torch.eye(n)` returns the `n × n` identity matrix.

`y = torch.eye(n, m)` returns an `n × m` identity matrix with ones on the diagonal and zeros elsewhere.

[res] torch.histc([res,] x [,nbins, min_value, max_value])

`y = torch.histc(x)` returns the histogram of the elements in `x`.

By default the elements are sorted into 100 equally spaced bins between the minimum and maximum values of `x`.

`y = torch.histc(x, n)` same as above with `n` bins.

`y = torch.histc(x, n, min, max)` same as above with `n` bins and `[min, max]` as elements range.

[res] torch.bhistc([res,] x [,nbins, min_value, max_value])

`y = torch.bhistc(x)` returns the histogram of the elements in 2d tensor `x` along the last dimension.

By default the elements are sorted into 100 equally spaced bins between the minimum and maximum values of `x`.

`y = torch.bhistc(x, n)` same as above with `n` bins.

`y = torch.bhistc(x, n, min, max)` same as above with `n` bins and `[min, max]` as elements range.

```
x = torch.Tensor(3, 6)

> x[1] = torch.Tensor{ 2, 4, 2, 2, 5, 4 }
> x[2] = torch.Tensor{ 3, 5, 1, 5, 3, 5 }
> x[3] = torch.Tensor{ 3, 4, 2, 5, 5, 1 }

> x
  2  4  2  2  5  4
  3  5  1  5  3  5
  3  4  2  5  5  1
[torch.DoubleTensor of size 3x6]

> torch.bhistc(x, 5, 1, 5)
  0  3  0  2  1
  1  0  2  0  3
  1  1  1  1  2
[torch.DoubleTensor of size 3x5]

> y = torch.Tensor(1, 6):copy(x[1])

> torch.bhistc(y, 5)
  3  0  2  0  1
[torch.DoubleTensor of size 1x5]
```

[res] torch.linspace([res,] x1, x2, [,n])

`y = torch.linspace(x1, x2)` returns a one-dimensional Tensor of size 100 equally

spaced points between `x1` and `x2`.

`y = torch.linspace(x1, x2, n)` returns a one-dimensional Tensor of `n` equally spaced points between `x1` and `x2`.

`[res] torch.logspace([res,] x1, x2, [,n])`

`y = torch.logspace(x1, x2)` returns a one-dimensional Tensor of 100 logarithmically equally spaced points between 10^{x1} and 10^{x2} .

`y = torch.logspace(x1, x2, n)` returns a one-dimensional Tensor of `n` logarithmically equally spaced points between 10^{x1} and 10^{x2} .

`[res] torch.multinomial([res,], p, n, [,replacement])`

`y = torch.multinomial(p, n)` returns a Tensor `y` where each row contains `n` indices sampled from the [multinomial probability distribution](#) located in the corresponding row of Tensor `p`.

The rows of `p` do not need to sum to one (in which case we use the values as weights), but must be non-negative and have a non-zero sum.

Indices are ordered from left to right according to when each was sampled (first samples are placed in first column).

If `p` is a vector, `y` is a vector size `n`.

If `p` is a `m`-rows matrix, `y` is an `m × n` matrix.

If `replacement` is `true`, samples are drawn **with replacement**.

If not, they are drawn **without replacement**, which means that when a sample index is drawn for a row, it cannot be drawn again for that row.

This implies the constraint that `n` must be lower than `p` length (or number of columns of `p` if it is a matrix).

The default value for `replacement` is `false`.

```
p = torch.Tensor{1, 1, 0.5, 0}
a = torch.multinomial(p, 10000, true)

> a
```

```
...
[torch.LongTensor of dimension 10000]

> for i = 1, 4 do print(a:eq(i):sum()) end
3967
4016
2017
0
```

Note: If you use the function with a given result `Tensor`, i.e. of the function prototype:
`torch.multinomial(res, p, n [, replacement])` then you will have to call it slightly differently as:

```
p.multinomial(res, p, n, replacement) -- p.multinomial instead of
torch.multinomial
```

This is due to the fact that the result here is of a `LongTensor` type, and we do not define a `torch.multinomial` over `long Tensor`s.

[res] torch.ones([res,] m [,n...])

`y = torch.ones(n)` returns a one-dimensional `Tensor` of size `n` filled with ones.

`y = torch.ones(m, n)` returns a `m × n` `Tensor` filled with ones.

For more than 4 dimensions, you can use a storage as argument: `y = torch.ones(torch.LongStorage{m, n, k, l, o})`.

[res] torch.rand([res,] [gen,] m [,n...])

`y = torch.rand(n)` returns a one-dimensional `Tensor` of size `n` filled with random numbers from a uniform distribution on the interval `[0, 1)`.

`y = torch.rand(m, n)` returns a `m × n` `Tensor` of random numbers from a uniform distribution on the interval `[0, 1)`.

For more than 4 dimensions, you can use a storage as argument: `y = torch.rand(torch.LongStorage{m, n, k, l, o})`.

`y = torch.rand(gen, m, n)` returns a `m × n` Tensor of random numbers from a uniform distribution on the interval `[0, 1)`, using a non-global random number generator `gen` created by `torch.Generator()`.

`[res] torch.randn([res,] [gen,] m [,n...])`

`y = torch.randn(n)` returns a one-dimensional Tensor of size `n` filled with random numbers from a normal distribution with mean zero and variance one.

`y = torch.randn(m, n)` returns a `m × n` Tensor of random numbers from a normal distribution with mean zero and variance one.

For more than 4 dimensions, you can use a storage as argument: `y = torch.randn(torch.LongStorage{m, n, k, l, o})`.

`y = torch.randn(gen, m, n)` returns a `m × n` Tensor of random numbers from a normal distribution with mean zero and variance one, using a non-global random number generator `gen` created by `torch.Generator()`.

`[res] torch.range([res,] x, y [,step])`

`y = torch.range(x, y)` returns a Tensor of size `floor((y - x) / step) + 1` with values from `x` to `y` with step `step` (default to 1).

```
> torch.range(2, 5)
2
3
4
5
[torch.DoubleTensor of size 4]

> torch.range(2, 5, 1.2)
2.0000
3.2000
4.4000
[torch.DoubleTensor of size 3]
```

`[res] torch.randperm([res,] [gen,] n)`

`y = torch.randperm(n)` returns a random permutation of integers from 1 to `n`.

`y = torch.randperm(gen, n)` returns a random permutation of integers from 1 to `n`, using a non-global random number generator `gen` created by [torch.Generator\(\)](#).

[res] torch.reshape([res,] x, m [,n...])

`y = torch.reshape(x, m, n)` returns a new `m × n` Tensor `y` whose elements are taken rowwise from `x`, which must have `m * n` elements. The elements are copied into the new Tensor.

For more than 4 dimensions, you can use a storage: `y = torch.reshape(x, torch.LongStorage{m, n, k, l, o})`.

[res] torch.tril([res,] x [,k])

`y = torch.tril(x)` returns the lower triangular part of `x`, the other elements of `y` are set to 0.

`torch.tril(x, k)` returns the elements on and below the `k`-th diagonal of `x` as non-zero. `k = 0` is the main diagonal, `k > 0` is above the main diagonal and `k < 0` is below the main diagonal.

[res] torch.triu([res,] x [,k])

`y = torch.triu(x)` returns the upper triangular part of `x`, the other elements of `y` are set to 0.

`torch.triu(x, k)` returns the elements on and above the `k`-th diagonal of `x` as non-zero. `k = 0` is the main diagonal, `k > 0` is above the main diagonal and `k < 0` is below the main diagonal.

[res] torch.zeros([res,] x)

`y = torch.zeros(n)` returns a one-dimensional Tensor of size `n` filled with zeros.

`y = torch.zeros(m, n)` returns a `m × n` `Tensor` filled with zeros.

For more than 4 dimensions, you can use a storage: `y = torch.zeros(torch.LongStorage{m, n, k, l, o})`.

Element-wise Mathematical Operations

`[res] torch.abs([res,] x)`

`y = torch.abs(x)` returns a new `Tensor` with the absolute values of the elements of `x`.

`x:abs()` replaces all elements in-place with the absolute values of the elements of `x`.

`[res] torch.sign([res,] x)`

`y = torch.sign(x)` returns a new `Tensor` with the sign (`+/- 1`) of the elements of `x`.

`x:sign()` replaces all elements in-place with the sign of the elements of `x`.

`[res] torch.acos([res,] x)`

`y = torch.acos(x)` returns a new `Tensor` with the arcosine of the elements of `x`.

`x:acos()` replaces all elements in-place with the arcosine of the elements of `x`.

`[res] torch.asin([res,] x)`

`y = torch.asin(x)` returns a new `Tensor` with the arcsine of the elements of `x`.

`x:asin()` replaces all elements in-place with the arcsine of the elements of `x`.

`[res] torch.atan([res,] x)`

`y = torch.atan(x)` returns a new `Tensor` with the arctangent of the elements of `x`.

`x:atan()` replaces all elements in-place with the arctangent of the elements of `x`.

`[res] torch.atan2([res,] x, y)`

`y = torch.atan2(x, y)` returns a new `Tensor` with the arctangent of the elements of `x` and `y`.

`x:atan2()` replaces all elements in-place with the arctangent of the elements of `x` and `y`.

`[res] torch.ceil([res,] x)`

`y = torch.ceil(x)` returns a new `Tensor` with the values of the elements of `x` rounded up to the nearest integers.

`x:ceil()` replaces all elements in-place with the values of the elements of `x` rounded up to the nearest integers.

`[res] torch.cos([res,] x)`

`y = torch.cos(x)` returns a new `Tensor` with the cosine of the elements of `x`.

`x:cos()` replaces all elements in-place with the cosine of the elements of `x`.

`[res] torch.cosh([res,] x)`

`y = torch.cosh(x)` returns a new `Tensor` with the hyperbolic cosine of the elements of `x`.

`x:cosh()` replaces all elements in-place with the hyperbolic cosine of the elements of `x`.

`[res] torch.exp([res,] x)`

`y = torch.exp(x)` returns, for each element in `x`, e (*Neper number*, the base of natural logarithms) raised to the power of the element in `x`.

`x:exp()` returns, for each element in `x`, e raised to the power of the element in `x`.

[res] torch.floor([res,] x)

`y = torch.floor(x)` returns a new `Tensor` with the values of the elements of `x` rounded down to the nearest integers.

`x:floor()` replaces all elements in-place with the values of the elements of `x` rounded down to the nearest integers.

[res] torch.log([res,] x)

`y = torch.log(x)` returns a new `Tensor` with the natural logarithm of the elements of `x`.

`x:log()` replaces all elements in-place with the natural logarithm of the elements of `x`.

[res] torch.log1p([res,] x)

`y = torch.log1p(x)` returns a new `Tensor` with the natural logarithm of the elements of `x + 1`.

`x:log1p()` replaces all elements in-place with the natural logarithm of the elements of `x + 1`.

This function is more accurate than `log` for small values of `x`.

x:neg()

`x:neg()` replaces all elements in-place with the sign-reversed values of the elements of `x`.

x:cinv()

`x:cinv()` replaces all elements in-place with $1.0 / x$.

[res] torch.pow([res,] x, n)

Let `x` be a `Tensor` and `n` a number.

`y = torch.pow(x, n)` returns a new `Tensor` with the elements of `x` to the power of `n`.

`y = torch.pow(n, x)` returns, a new `Tensor` with `n` to the power of the elements of `x`.

`x:pow(n)` replaces all elements in-place with the elements of `x` to the power of `n`.

`torch.pow(x, n, x)` replaces all elements in-place with `n` to the power of the elements of `x`.

[res] torch.round([res,] x)

`y = torch.round(x)` returns a new `Tensor` with the values of the elements of `x` rounded to the nearest integers.

`x:round()` replaces all elements in-place with the values of the elements of `x` rounded to the nearest integers.

[res] torch.sin([res,] x)

`y = torch.sin(x)` returns a new `Tensor` with the sine of the elements of `x`.

`x:sin()` replaces all elements in-place with the sine of the elements of `x`.

[res] torch.sinh([res,] x)

`y = torch.sinh(x)` returns a new `Tensor` with the hyperbolic sine of the elements of `x`.

`x:sinh()` replaces all elements in-place with the hyperbolic sine of the elements of `x`.

[res] torch.sqrt([res,] x)

`y = torch.sqrt(x)` returns a new `Tensor` with the square root of the elements of `x`.

`x: sqrt()` replaces all elements in-place with the square root of the elements of `x`.

[res] torch.rsqrt([res,] x)

`y = torch.rsqrt(x)` returns a new `Tensor` with the reciprocal of the square root of the elements of `x`.

`x: rsqrt()` replaces all elements in-place with the reciprocal of the square root of the elements of `x`.

[res] torch.tan([res,] x)

`y = torch.tan(x)` returns a new `Tensor` with the tangent of the elements of `x`.

`x: tan()` replaces all elements in-place with the tangent of the elements of `x`.

[res] torch.tanh([res,] x)

`y = torch.tanh(x)` returns a new `Tensor` with the hyperbolic tangent of the elements of `x`.

`x: tanh()` replaces all elements in-place with the hyperbolic tangent of the elements of `x`.

[res] torch.sigmoid([res,] x)

`y = torch.sigmoid(x)` returns a new `Tensor` with the sigmoid of the elements of `x`.

`x: sigmoid()` replaces all elements in-place with the sigmoid of the elements of `x`.

[res] torch.trunc([res,] x)

`y = torch.trunc(x)` returns a new `Tensor` with the truncated integer values of the elements of `x`.

`x:trunc()` replaces all elements in-place with the truncated integer values of the elements of `x`.

[res] torch.frac([res,] x)

`y = torch.frac(x)` returns a new `Tensor` with the fractional portion of the elements of `x`.

`x:frac()` replaces all elements in-place with the fractional portion of the elements of `x`.

Basic operations

In this section, we explain basic mathematical operations for `Tensor` s.

[boolean] equal([tensor1,] tensor2)

Returns `true` iff the dimensions and values of `tensor1` and `tensor2` are exactly the same.

```
x = torch.Tensor{1,2,3}
y = torch.Tensor{1,2,3}
> x:equal(y)
true

y = torch.Tensor{1,2,4}
> x:equal(y)
false
```

Note that `a:equal(b)` is more efficient than `a:eq(b):all()` as it avoids allocation of a temporary tensor and can short-circuit.

[res] torch.add([res,] tensor, value)

Add the given value to all elements in the `Tensor` .

`y = torch.add(x, value)` returns a new `Tensor` .

`x:add(value)` add `value` to all elements in place.

`[res] torch.add([res,] tensor1, tensor2)`

Add `tensor1` to `tensor2` and put result into `res` .

The number of elements must match, but sizes do not matter.

```
> x = torch.Tensor(2, 2):fill(2)
> y = torch.Tensor(4):fill(3)
> x:add(y)
> x
  5  5
  5  5
[torch.DoubleTensor of size 2x2]
```

`y = torch.add(a, b)` returns a new `Tensor` .

`torch.add(y, a, b)` puts `a + b` in `y` .

`a:add(b)` accumulates all elements of `b` into `a` .

`y:add(a, b)` puts `a + b` in `y` .

`[res] torch.add([res,] tensor1, value, tensor2)`

Multiply elements of `tensor2` by the scalar `value` and add it to `tensor1` .

The number of elements must match, but sizes do not matter.

```
> x = torch.Tensor(2, 2):fill(2)
> y = torch.Tensor(4):fill(3)
> x:add(2, y)
> x
  8  8
  8  8
[torch.DoubleTensor of size 2x2]
```


`x:add(value, y)` multiply-accumulates values of `y` into `x`.

`z:add(x, value, y)` puts the result of `x + value * y` in `z`.

`torch.add(x, value, y)` returns a new `Tensor` `x + value * y`.

`torch.add(z, x, value, y)` puts the result of `x + value * y` in `z`.

`tensor:csub(value)`

Subtracts the given value from all elements in the `Tensor`, in place.

`tensor:csub(tensor2)`

Subtracts `tensor2` from `tensor`, in place.

The number of elements must match, but sizes do not matter.

```
> x = torch.Tensor(2, 2):fill(8)
> y = torch.Tensor(4):fill(3)
> x:csub(y)
> x
  5  5
  5  5
[torch.DoubleTensor of size 2x2]
```

`a:csub(b)` put `a - b` into `a`.

`[res] torch.mul([res,] tensor1, value)`

Multiply all elements in the `Tensor` by the given `value`.

`z = torch.mul(x, 2)` will return a new `Tensor` with the result of `x * 2`.

`torch.mul(z, x, 2)` will put the result of `x * 2` in `z`.

`x:mul(2)` will multiply all elements of `x` with `2` in-place.

`z:mul(x, 2)` will put the result of `x * 2` in `z`.

[res] torch.clamp([res,] tensor, min_value, max_value)

Clamp all elements in the `Tensor` into the range `[min_value, max_value]` . ie:

$$y_i = \begin{cases} \text{min_value, if } x_i < \text{min_value} \\ x_i, & \text{if } \text{min_value} \leq x_i \leq \text{max_value} \\ \text{max_value, if } x_i > \text{max_value} \end{cases}$$

`z = torch.clamp(x, 0, 1)` will return a new `Tensor` with the result of `x` bounded between `0` and `1` .

`torch.clamp(z, x, 0, 1)` will put the result in `z` .

`x:clamp(0, 1)` will perform the clamp operation in place (putting the result in `x`).

`z:clamp(x, 0, 1)` will put the result in `z` .

[res] torch.cmul([res,] tensor1, tensor2)

Element-wise multiplication of `tensor1` by `tensor2` .

The number of elements must match, but sizes do not matter.

```
> x = torch.Tensor(2, 2):fill(2)
> y = torch.Tensor(4):fill(3)
> x:cmul(y)
> = x
  6  6
  6  6
[torch.DoubleTensor of size 2x2]
```

`z = torch.cmul(x, y)` returns a new `Tensor` .

`torch.cmul(z, x, y)` puts the result in `z` .

`y:cmul(x)` multiplies all elements of `y` with corresponding elements of `x` .

`z:cmul(x, y)` puts the result in `z` .

[res] torch.cpow([res,] tensor1, tensor2)

Element-wise power operation, taking the elements of `tensor1` to the powers given by elements of `tensor2`.

The number of elements must match, but sizes do not matter.

```
> x = torch.Tensor(2, 2):fill(2)
> y = torch.Tensor(4):fill(3)
> x:cpow(y)
> x
  8  8
  8  8
[torch.DoubleTensor of size 2x2]
```

`z = torch.cpow(x, y)` returns a new `Tensor`.

`torch.cpow(z, x, y)` puts the result in `z`.

`y:cpow(x)` takes all elements of `y` to the powers given by the corresponding elements of `x`.

`z:cpow(x, y)` puts the result in `z`.

[res] torch.addcmul([res,] x [,value], tensor1, tensor2)

Performs the element-wise multiplication of `tensor1` by `tensor2`, multiply the result by the scalar `value` (1 if not present) and add it to `x`.

The number of elements must match, but sizes do not matter.

```
> x = torch.Tensor(2, 2):fill(2)
> y = torch.Tensor(4):fill(3)
> z = torch.Tensor(2, 2):fill(5)
> x:addcmul(2, y, z)
> x
 32 32
 32 32
[torch.DoubleTensor of size 2x2]
```

`z:addcmul(value, x, y)` accumulates the result in `z`.

`torch.addcmul(z, value, x, y)` returns a new `Tensor` with the result.

`torch.addcmul(z, z, value, x, y)` puts the result in `z`.

[res] torch.div([res,] tensor, value)

Divide all elements in the `Tensor` by the given `value`.

`z = torch.div(x, 2)` will return a new `Tensor` with the result of `x / 2`.

`torch.div(z, x, 2)` will put the result of `x / 2` in `z`.

`x:div(2)` will divide all elements of `x` with `2` in-place.

`z:div(x, 2)` puts the result of `x / 2` in `z`.

[res] torch.cdiv([res,] tensor1, tensor2)

Performs the element-wise division of `tensor1` by `tensor2`.
The number of elements must match, but sizes do not matter.

```
> x = torch.Tensor(2, 2):fill(1)
> y = torch.range(1, 4)
> x:cdiv(y)
> x
  1.0000  0.5000
  0.3333  0.2500
[torch.DoubleTensor of size 2x2]
```

`z = torch.cdiv(x, y)` returns a new `Tensor`.

`torch.cdiv(z, x, y)` puts the result in `z`.

`y:cdiv(x)` divides all elements of `y` with corresponding elements of `x`.

`z:cdiv(x, y)` puts the result in `z`.

[res] torch.addcdiv([res,] x [,value], tensor1, tensor2)

Performs the element-wise division of `tensor1` by `tensor2`, multiply the result by the scalar `value` and add it to `x`.

The number of elements must match, but sizes do not matter.

```
> x = torch.Tensor(2, 2):fill(1)
> y = torch.range(1, 4)
> z = torch.Tensor(2, 2):fill(5)
> x:addcddiv(2, y, z)
> x
  1.4000  1.8000
  2.2000  2.6000
[torch.DoubleTensor of size 2x2]
```

`z:addcddiv(value, x, y)` accumulates the result in `z`.

`torch.addcddiv(z, value, x, y)` returns a new `Tensor` with the result.

`torch.addcddiv(z, z, value, x, y)` puts the result in `z`.

[res] torch.fmod([res,] tensor, value)

Computes remainder of division (rounded towards zero) of all elements in the `Tensor` by `value`.

This works both for integer and floating point numbers. It behaves the same as Lua built-in function `math.fmod()` and a little bit different from `torch remainder()` and `%` operator. For example:

```
> x = torch.Tensor({-3, 3})
> torch.fmod(x, 2)
-1
 1
[torch.DoubleTensor of size 2]

> torch.fmod(x, -2)
-1
 1
[torch.DoubleTensor of size 2]

> torch.remainder(x, 2)
 1
 1
```

```
[torch.DoubleTensor of size 2]

> torch.remainder(x, -2)
-1
-1
[torch.DoubleTensor of size 2]
```

`z = torch.fmod(x, 2)` will return a new `Tensor` with the result of `math.fmod(x, 2)`.

`torch.fmod(z, x, 2)` will put the result of `math.fmod(x, 2)` in `z`.

`x:fmod(2)` will replace all elements of `x` the result of `math.fmod(x, 2)` in-place.

`z:fmod(x, 2)` puts the result of `math.fmod(x, 2)` in `z`.

[res] torch.remainder([res,] tensor, value)

Computes remainder of division (rounded to nearest) of all elements in the `Tensor` by `value`. This works both for integer and floating point numbers. It behaves the same as `%` operator and can be expressed as $a \% b = a - b * \text{floor}(a/b)$. See `torch.fmod()` for comparison.

`z = torch.remainder(x, 2)` will return a new `Tensor` with the result of `x % 2`.

`torch.remainder(z, x, 2)` will put the result of `x % 2` in `z`.

`x:remainder(2)` will replace all elements of `x` the result of `x % 2` in-place.

`z:remainder(x, 2)` puts the result of `x % 2` in `z`.

[res] torch.mod([res,] tensor, value)

This function is deprecated and exists only for compatibility with previous versions. Please use `torch.fmod()` or `torch.remainder()` instead.

[res] torch.cfmod([res,] tensor1, tensor2)

Computes the element-wise remainder of the division (rounded towards zero) of `tensor1` by `tensor2`.

The number of elements must match, but sizes do not matter.

```

> x = torch.Tensor([[3, 3], [-3, -3]])
> y = torch.Tensor([[2, -2], {2, -2}])
> x:cfmod(y)
  1  1
 -1 -1
[torch.DoubleTensor of size 2x2]

```

`z = torch.cfmod(x, y)` returns a new `Tensor` .

`torch.cfmod(z, x, y)` puts the result in `z` .

`y:cfmod(x)` replaces all elements of `y` by their remainders of division (rounded towards zero) by corresponding elements of `x` .

`z:cfmod(x, y)` puts the result in `z` .

[res] torch.cremainder([res,] tensor1, tensor2)

Computes element-wise remainder of the division (rounded to nearest) of `tensor1` by `tensor2` .

The number of elements must match, but sizes do not matter.

```

> x = torch.Tensor([[3, 3], [-3, -3]])
> y = torch.Tensor([[2, -2], {2, -2}])
> x:cfmod(y)
  1  1
 -1 -1
[torch.DoubleTensor of size 2x2]

```

`z = torch.cremainder(x, y)` returns a new `Tensor` .

`torch.cremainder(z, x, y)` puts the result in `z` .

`y:cremainder(x)` replaces all elements of `y` by their remainders of division (rounded to nearest) by corresponding elements of `x` .

`z:cremainder(x, y)` puts the result in `z` .

[res] torch.cmod([res,] tensor1, tensor2)

This function is deprecated and exists only for compatibility with previous versions. Please use `torch.cfmmod()` or `torch.cremainder()` instead.

[number] torch.dot(tensor1, tensor2)

Performs the dot product between `tensor1` and `tensor2`.

The number of elements must match: both `Tensor` s are seen as a 1D vector.

```
> x = torch.Tensor(2, 2):fill(2)
> y = torch.Tensor(4):fill(3)
> x:dot(y)
24
```

`torch.dot(x, y)` returns dot product of `x` and `y`.

`x:dot(y)` returns dot product of `x` and `y`.

[res] torch.addmv([res,] [v1,] vec1, [v2,] mat, vec2)

Performs a matrix-vector multiplication between `mat` (2D `Tensor`) and `vec2` (1D `Tensor`) and add it to `vec1`.

Optional values `v1` and `v2` are scalars that multiply `vec1` and `vec2` respectively.

In other words,

```
res = (v1 * vec1) + (v2 * (mat * vec2))
```

Sizes must respect the matrix-multiplication operation: if `mat` is a `n × m` matrix, `vec2` must be vector of size `m` and `vec1` must be a vector of size `n`.

```
> x = torch.Tensor(3):fill(0)
> M = torch.Tensor(3, 2):fill(3)
> y = torch.Tensor(2):fill(2)
> x:addmv(M, y)
```



```
> x
12
12
12
[torch.DoubleTensor of size 3]
```

`torch.addmv(x, y, z)` returns a new `Tensor` with the result.

`torch.addmv(r, x, y, z)` puts the result in `r`.

Differences when used as a method

`x:addmv(y, z)` does $x = x + y * z$

`r:addmv(x, y, z)` does $r = x + y * z$ if `x` is a vector

`r:addmv(s, y, z)` does $r = r + s * y * z$ if `s` is a scalar.

`r:addmv(x, s, y, z)` does $r = x + s * y * z$ if `s` is a scalar and `x` is a vector.

`r:addmv(s1, s2, y, z)` does $r = s1 * r + s2 * y * z$ if `s1` and `s2` are scalars.

The last example does not accurately fit into the function signature, and needs a special mention. It changes the function signature to:

```
[vec1] = vec1:addmv([v1,] [v2,] mat, vec2)
```

[res] torch.addr([res,] [v1,] mat, [v2,] vec1, vec2)

Performs the outer-product between `vec1` (1D `Tensor`) and `vec2` (1D `Tensor`).

Optional values `v1` and `v2` are scalars that multiply `mat` and `vec1` [out] `vec2` respectively.

In other words,

```
res_ij = (v1 * mat_ij) + (v2 * vec1_i * vec2_j)
```

If `vec1` is a vector of size `n` and `vec2` is a vector of size `m`, then `mat` must be a matrix of size `n × m`.

```
> x = torch.range(1, 3)
```

```

> y = torch.range(1, 2)
> M = torch.Tensor(3, 2):zero()
> M:addr(x, y)
  1  2      --      | 0 0|      | 1 2|
  2  4      -- = 1*| 0 0| + 1*| 2 4|
  3  6      --      | 0 0|      | 3 6|
[torch.DoubleTensor of size 3x2]
-- default values of v1 and v2 are 1.

> M:addr(2, 1, x, y)
  3  6      --      | 1 2|      | 1 2|
  6 12      -- = 2*| 2 4| + 1*| 2 4|
  9 18      --      | 3 6|      | 3 6|
[torch.DoubleTensor of size 3x2]

> A = torch.range(1, 6):resize(3, 2)
> A
  1  2
  3  4
  5  6
[torch.DoubleTensor of size 3x2]
> M:addr(2, A, 1, x, y)
  3  6      --      | 1 2|      | 1 2|
  8 12      -- = 2*| 3 4| + 1*| 2 4|
 13 18      --      | 5 6|      | 3 6|
[torch.DoubleTensor of size 3x2]

```

`torch.addr(M, x, y)` returns the result in a new `Tensor` .

`torch.addr(r, M, x, y)` puts the result in `r` .

`M:addr(x, y)` puts the result in `M` .

`r:addr(M, x, y)` puts the result in `r` .

`[res] torch.addmm([res,] [v1,] M, [v2,] mat1, mat2)`

Performs a matrix-matrix multiplication between `mat1` (2D `Tensor`) and `mat2` (2D `Tensor`).

Optional values `v1` and `v2` are scalars that multiply `M` and `mat1 * mat2` respectively.

In other words,

```
res = (v1 * M) + (v2 * mat1 * mat2)
```

If `mat1` is a $n \times m$ matrix, `mat2` a $m \times p$ matrix, `M` must be a $n \times p$ matrix.

`torch.addmm(M, mat1, mat2)` returns the result in a new `Tensor`.

`torch.addmm(r, M, mat1, mat2)` puts the result in `r`.

Differences when used as a method

`M:addmm(mat1, mat2)` does $M = M + mat1 * mat2$.

`r:addmm(M, mat1, mat2)` does $r = M + mat1 * mat2$.

`r:addmm(v1, M, v2, mat1, mat2)` does $r = (v1 * M) + (v2 * mat1 * mat2)$.

`M:addmm(v1, v2, mat1, mat2)` does $M = (v1 * M) + (v2 * mat1 * mat2)$.

The last example does not accurately fit into the function signature, and needs a special mention. It changes the function signature to:

```
[M] = M:addmm([v1,] [v2,] mat1, mat2)
```

[res] torch.addbmm([res,] [v1,] M, [v2,] batch1, batch2)

Batch matrix matrix product of matrices stored in `batch1` and `batch2`, with a reduced add step (all matrix multiplications get accumulated in a single place).

`batch1` and `batch2` must be 3D `Tensor`s each containing the same number of matrices. If `batch1` is a $b \times n \times m$ `Tensor`, `batch2` a $b \times m \times p$ `Tensor`, `res` will be a $n \times p$ `Tensor`.

In other words,

```
res = (v1 * M) + (v2 * sum(batch1_i * batch2_i, i = 1, b))
```

`torch.addbmm(M, x, y)` puts the result in a new `Tensor`.

`M:addbmm(x, y)` puts the result in `M`, resizing `M` if necessary.

`M:addbmm(beta, M2, alpha, x, y)` puts the result in `M`, resizing `M` if necessary.

[res] torch.baddbmm([res,] [v1,] M, [v2,] batch1, batch2)

Batch matrix matrix product of matrices stored in `batch1` and `batch2`, with batch add.

`batch1` and `batch2` must be 3D `Tensor`s each containing the same number of matrices. If `batch1` is a $b \times n \times m$ `Tensor`, `batch2` a $b \times m \times p$ `Tensor`, `res` will be a $b \times n \times p$ `Tensor`.

In other words,

$$\text{res}_i = (v1 * M_i) + (v2 * \text{batch1}_i * \text{batch2}_i)$$

`torch.baddbmm(M, x, y)` puts the result in a new `Tensor`.

`M:baddbmm(x, y)` puts the result in `M`, resizing `M` if necessary.

`M:baddbmm(beta, M2, alpha, x, y)` puts the result in `M`, resizing `M` if necessary.

[res] torch.mv([res,] mat, vec)

Matrix vector product of `mat` and `vec`.

Sizes must respect the matrix-multiplication operation: if `mat` is a $n \times m$ matrix, `vec` must be vector of size `m` and `res` must be a vector of size `n`.

`torch.mv(x, y)` puts the result in a new `Tensor`.

`torch.mv(M, x, y)` puts the result in `M`.

`M:mv(x, y)` puts the result in `M`.

[res] torch.mm([res,] mat1, mat2)

Matrix matrix product of `mat1` and `mat2`.

If `mat1` is a $n \times m$ matrix, `mat2` a $m \times p$ matrix, `res` must be a $n \times p$ matrix.

`torch.mm(x, y)` puts the result in a new `Tensor`.

`torch.mm(M, x, y)` puts the result in `M`.

`M:mm(x, y)` puts the result in `M`.

`[res] torch.bmm([res,] batch1, batch2)`

Batch matrix matrix product of matrices stored in `batch1` and `batch2`.

`batch1` and `batch2` must be 3D Tensor s each containing the same number of matrices.

If `batch1` is a $b \times n \times m$ Tensor , `batch2` a $b \times m \times p$ Tensor , `res` will be a $b \times n \times p$ Tensor .

`torch.bmm(x, y)` puts the result in a new Tensor .

`torch.bmm(M, x, y)` puts the result in `M`, resizing `M` if necessary.

`M:bmm(x, y)` puts the result in `M`, resizing `M` if necessary.

`[res] torch.ger([res,] vec1, vec2)`

Outer product of `vec1` and `vec2`.

If `vec1` is a vector of size `n` and `vec2` is a vector of size `m`, then `res` must be a matrix of size $n \times m$.

`torch.ger(x, y)` puts the result in a new Tensor .

`torch.ger(M, x, y)` puts the result in `M`.

`M:ger(x, y)` puts the result in `M`.

`[res] torch.lerp([res,] a, b, weight)`

Linear interpolation of two scalars or tensors based on a weight: `res = a + weight * (b - a)`

`torch.lerp(a, b, weight)` puts the result in a new Tensor if `a` and `b` are tensors. If `a` and `b` are scalars the functions returns a number.

`torch.lerp(M, a, b, weight)` puts the result in `M`.

`M:lerp(a, b, weight)` puts the result in `M`.

Overloaded operators

It is possible to use basic mathematical operators like `+`, `-`, `/`, `*` and `%` with `Tensor` s. These operators are provided as a convenience.

While they might be handy, they create and return a new `Tensor` containing the results. They are thus not as fast as the operations available in the [previous section](#).

Another important point to note is that these operators are only overloaded when the first operand is a `Tensor` .

For example, this will NOT work:

```
> x = 5 + torch.rand(3)
```

Addition and subtraction

You can add a `Tensor` to another one with the `+` operator.

Subtraction is done with `-` .

The number of elements in the `Tensor` s must match, but the sizes do not matter.

The size of the returned `Tensor` will be the size of the first `Tensor` .

```
> x = torch.Tensor(2, 2):fill(2)
> y = torch.Tensor(4):fill(3)
> z = x + y
 5  5
 5  5
[torch.DoubleTensor of size 2x2]

> w = y - x
 1
 1
 1
 1
[torch.DoubleTensor of size 4]
```

A scalar might also be added or subtracted to a `Tensor` .

The scalar needs to be on the right of the operator.

```
> x = torch.Tensor(2, 2):fill(2)
> = x + 3
5 5
5 5
[torch.DoubleTensor of size 2x2]
```

Negation

A `Tensor` can be negated with the `-` operator placed in front:

```
> x = torch.Tensor(2, 2):fill(2)
> = -x
-2 -2
-2 -2
[torch.DoubleTensor of size 2x2]
```

Multiplication

Multiplication between two `Tensor`s is supported with the `*` operators. The result of the multiplication depends on the sizes of the `Tensor`s.

- 1D and 1D: Returns the dot product between the two `Tensor`s (scalar).
- 2D and 1D: Returns the matrix-vector operation between the two `Tensor`s (1D `Tensor`).
- 2D and 2D: Returns the matrix-matrix operation between the two `Tensor`s (2D `Tensor`).

Sizes must be conformant for the corresponding operation.

A `Tensor` might also be multiplied by a scalar. The scalar might be on the right or left of the operator.

Examples:

```
> M = torch.Tensor(2, 2):fill(2)
> N = torch.Tensor(2, 4):fill(3)
> x = torch.Tensor(2):fill(4)
```

```

> y = torch.Tensor(2):fill(5)
> = x * y -- dot product
40

> = M * x --- matrix-vector
16
16
[torch.DoubleTensor of size 2]

> = M * N -- matrix-matrix
12 12 12 12
12 12 12 12
[torch.DoubleTensor of size 2x4]

```

Division and Modulo (remainder)

Only the division of a `Tensor` by a scalar is supported with the operator `/`.

Example:

```

> x = torch.Tensor(2, 2):fill(2)
> = x/3
0.6667 0.6667
0.6667 0.6667
[torch.DoubleTensor of size 2x2]

```

Similarly, the remainder of the division of a `Tensor`'s elements by a scalar can be obtained with the operator `%`.

Example:

```

x = torch.Tensor({{1,2},{3,4}})
= x % 3
1 2
0 1
[torch.Tensor of size 2x2]

```


Column or row-wise operations (dimension-wise operations)

[res] torch.cross([res,] a, b [,n])

`y = torch.cross(a, b)` returns the cross product of `a` and `b` along the first dimension of length 3.

`y = torch.cross(a, b, n)` returns the cross product of vectors in dimension `n` of `a` and `b`.

`a` and `b` must have the same size, and both `a:size(n)` and `b:size(n)` must be 3.

[res] torch.cumprod([res,] x [,dim])

`y = torch.cumprod(x)` returns the cumulative product of the elements of `x`, performing the operation over the last dimension.

`y = torch.cumprod(x, n)` returns the cumulative product of the elements of `x`, performing the operation over dimension `n`.

```
-- 1. cumulative product for a vector
> A = torch.range(1, 5)
> A
  1
  2
  3
  4
  5
[torch.DoubleTensor of size 5]

> B = torch.cumprod(A)
> B
  1      -- B(1) = A(1) = 1
  2      -- B(2) = A(1)*A(2) = 1*2 = 2
  6      -- B(3) = A(1)*A(2)*A(3) = 1*2*3 = 6
 24      -- B(4) = A(1)*A(2)*A(3)*A(4) = 1*2*3*4 = 24
120      -- B(5) = A(1)*A(2)*A(3)*A(4)*A(5) = 1*2*3*4*5 = 120
```

```
[torch.DoubleTensor of size 5]
```

```
-- 2. cumulative product for a matrix
```

```
> A = torch.LongTensor({1, 4, 7}, {2, 5, 8}, {3, 6, 9})
```

```
> A
```

```
1  4  7
2  5  8
3  6  9
```

```
[torch.LongTensor of size 3x3]
```

```
> B = torch.cumprod(A)
```

```
> B
```

```
1    4    7
2   20   56
6  120  504
```

```
[torch.LongTensor of size 3x3]
```

```
-- Why?
```

```
-- B(1, 1) = A(1, 1) = 1
```

```
-- B(2, 1) = A(1, 1)*A(2, 1) = 1*2 = 2
```

```
-- B(3, 1) = A(1, 1)*A(2, 1)*A(3, 1) = 1*2*3 = 6
```

```
-- B(1, 2) = A(1, 2) = 4
```

```
-- B(2, 2) = A(1, 2)*A(2, 2) = 4*5 = 20
```

```
-- B(3, 2) = A(1, 2)*A(2, 2)*A(3, 2) = 4*5*6 = 120
```

```
-- B(1, 3) = A(1, 3) = 7
```

```
-- B(2, 3) = A(1, 3)*A(2, 3) = 7*8 = 56
```

```
-- B(3, 3) = A(1, 3)*A(2, 3)*A(3, 3) = 7*8*9 = 504
```

```
-- 3. cumulative product along 2-dim
```

```
> B = torch.cumprod(A, 2)
```

```
> B
```

```
1    4   28
2   10   80
3   18  162
```

```
[torch.LongTensor of size 3x3]
```

```
-- Why?
```

```
-- B(1, 1) = A(1, 1) = 1
```

```
-- B(1, 2) = A(1, 1)*A(1, 2) = 1*4 = 4
```

```
-- B(1, 3) = A(1, 1)*A(1, 2)*A(1, 3) = 1*4*7 = 28
```

```
-- B(2, 1) = A(2, 1) = 2
```

```
-- B(2, 2) = A(2, 1)*A(2, 2) = 2*5 = 10
```

```
-- B(2, 3) = A(2, 1)*A(2, 2)*A(2, 3) = 2*5*8 = 80
```

```
-- B(3, 1) = A(3, 1) = 3
```

```
-- B(3, 2) = A(3, 1)*A(2, 3) = 3*6 = 18
```

```
-- B(3, 3) = A(3, 1)*A(2, 3)*A(3, 3) = 3*6*9 = 162
```

[res] torch.cumsum([res,] x [,dim])

`y = torch.cumsum(x)` returns the cumulative sum of the elements of `x`, performing the operation over the first dimension.

`y = torch.cumsum(x, n)` returns the cumulative sum of the elements of `x`, performing the operation over dimension `n`.

torch.max([resval, resind,] x [,dim])

`y = torch.max(x)` returns the single largest element of `x`.

`y, i = torch.max(x, 1)` returns the largest element in each column (across rows) of `x`, and a Tensor `i` of their corresponding indices in `x`.

`y, i = torch.max(x, 2)` performs the max operation for each row.

`y, i = torch.max(x, n)` performs the max operation over the dimension `n`.

```
> x = torch.randn(3, 3)
> x
 1.1994 -0.6290  0.6888
-0.0038 -0.0908 -0.2075
 0.3437 -0.9948  0.1216
[torch.DoubleTensor of size 3x3]

> torch.max(x)
1.1993977428735

> torch.max(x, 1)
 1.1994 -0.0908  0.6888
[torch.DoubleTensor of size 1x3]

 1  2  1
[torch.LongTensor of size 1x3]

> torch.max(x, 2)
1.1994
```

```
-0.0038
0.3437
[torch.DoubleTensor of size 3x1]

1
1
1
[torch.LongTensor of size 3x1]
```

[res] torch.mean([res,] x [,dim])

`y = torch.mean(x)` returns the mean of all elements of `x`.

`y = torch.mean(x, 1)` returns a `Tensor` `y` of the mean of the elements in each column of `x`.

`y = torch.mean(x, 2)` performs the mean operation for each row.

`y = torch.mean(x, n)` performs the mean operation over the dimension `n`.

torch.min([resval, resind,] x [,dim])

`y = torch.min(x)` returns the single smallest element of `x`.

`y, i = torch.min(x, 1)` returns the smallest element in each column (across rows) of `x`, and a `Tensor` `i` of their corresponding indices in `x`.

`y, i = torch.min(x, 2)` performs the min operation for each row.

`y, i = torch.min(x, n)` performs the min operation over the dimension `n`.

[res] torch.cmax([res,] tensor1, tensor2)

Compute the maximum of each pair of values in `tensor1` and `tensor2`.

`c = torch.cmax(a, b)` returns a new `Tensor` containing the element-wise maximum of `a` and `b`.

`a:cmax(b)` stores the element-wise maximum of `a` and `b` in `a`.

`c:cmax(a, b)` stores the element-wise maximum of `a` and `b` in `c`.

```
> a = torch.Tensor{1, 2, 3}
> b = torch.Tensor{3, 2, 1}
> torch.cmax(a, b)
3
2
3
[torch.DoubleTensor of size 3]
```

[res] torch.cmax([res,] tensor, value)

Compute the maximum between each value in `tensor` and `value`.

`c = torch.cmax(a, v)` returns a new `Tensor` containing the maxima of each element in `a` and `v`.

`a:cmax(v)` stores the maxima of each element in `a` and `v` in `a`.

`c:cmax(a, v)` stores the maxima of each element in `a` and `v` in `c`.

```
> a = torch.Tensor{1, 2, 3}
> torch.cmax(a, 2)
2
2
3
[torch.DoubleTensor of size 3]
```

[res] torch.cmin([res,] tensor1, tensor2)

Compute the minimum of each pair of values in `tensor1` and `tensor2`.

`c = torch.cmin(a, b)` returns a new `Tensor` containing the element-wise minimum of `a` and `b`.

`a:cmin(b)` stores the element-wise minimum of `a` and `b` in `a`.

`c:cmin(a, b)` stores the element-wise minimum of `a` and `b` in `c`.

```

> a = torch.Tensor{1, 2, 3}
> b = torch.Tensor{3, 2, 1}
> torch.cmin(a, b)
1
2
1
[torch.DoubleTensor of size 3]

```

torch.cmin([res,] tensor, value)

Compute the minimum between each value in `tensor` and `value`.

`c = torch.cmin(a, v)` returns a new `Tensor` containing the minima of each element in `a` and `v`.

`a:cmin(v)` stores the minima of each element in `a` and `v` in `a`.

`c:cmin(a, v)` stores the minima of each element in `a` and `v` in `c`.

```

> a = torch.Tensor{1, 2, 3}
> torch.cmin(a, 2)
1
2
2
[torch.DoubleTensor of size 3]

```

torch.median([resval, resind,] x [,dim])

`y = torch.median(x)` performs the median operation over the last dimension of `x` (one-before-middle in the case of an even number of elements).

`y, i = torch.median(x, 1)` returns the median element in each column (across rows) of `x`, and a `Tensor` `i` of their corresponding indices in `x`.

`y, i = torch.median(x, 2)` performs the median operation for each row.

`y, i = torch.median(x, n)` performs the median operation over the dimension `n`.

```
> x = torch.randn(3, 3)
> x
 0.7860  0.7687 -0.9362
 0.0411  0.5407 -0.3616
-0.0129 -0.2499 -0.5786
[torch.DoubleTensor of size 3x3]

> y, i = torch.median(x)
> y
 0.7687
 0.0411
-0.2499
[torch.DoubleTensor of size 3x1]

> i
 2
 1
 2
[torch.LongTensor of size 3x1]

> y, i = torch.median(x, 1)
> y
 0.0411  0.5407 -0.5786
[torch.DoubleTensor of size 1x3]

> i
 2  2  3
[torch.LongTensor of size 1x3]

> y, i = torch.median(x, 2)
> y
 0.7687
 0.0411
-0.2499
[torch.DoubleTensor of size 3x1]

> i
 2
 1
 2
[torch.LongTensor of size 3x1]
```

`torch.mode([resval, resind,] x [,dim])`

`y = torch.mode(x)` returns the most frequent element of `x` over its last dimension.

`y, i = torch.mode(x, 1)` returns the mode element in each column (across rows) of `x`, and a `Tensor` `i` of their corresponding indices in `x`.

`y, i = torch.mode(x, 2)` performs the mode operation for each row.

`y, i = torch.mode(x, n)` performs the mode operation over the dimension `n`.

`torch.kthvalue([resval, resind,] x, k [,dim])`

`y = torch.kthvalue(x, k)` returns the `k`-th smallest element of `x` over its last dimension.

`y, i = torch.kthvalue(x, k, 1)` returns the `k`-th smallest element in each column (across rows) of `x`, and a `Tensor` `i` of their corresponding indices in `x`.

`y, i = torch.kthvalue(x, k, 2)` performs the `k`-th value operation for each row.

`y, i = torch.kthvalue(x, k, n)` performs the `k`-th value operation over the dimension `n`.

`[res] torch.prod([res,] x [,n])`

`y = torch.prod(x)` returns the product of all elements in `x`.

`y = torch.prod(x, n)` returns a `Tensor` `y` whose size in dimension `n` is 1 and where elements are the product of elements of `x` with respect to dimension `n`.

```
> a = torch.Tensor{{{1, 2}, {3, 4}}, {{5, 6}, {7, 8}}}  
> a  
(1,.,.) =  
  1  2  
  3  4  
  
(2,.,.) =  
  5  6
```



```

 7  8
[torch.DoubleTensor of dimension 2x2x2]

> torch.prod(a, 1)
(1,.,.) =
  5  12
 21  32
[torch.DoubleTensor of dimension 1x2x2]

> torch.prod(a, 2)
(1,.,.) =
  3  8

(2,.,.) =
 35  48
[torch.DoubleTensor of size 2x1x2]

> torch.prod(a, 3)
(1,.,.) =
  2
 12

(2,.,.) =
 30
 56
[torch.DoubleTensor of size 2x2x1]

```

torch.sort([resval, resind,] x [,d] [,flag])

`y, i = torch.sort(x)` returns a Tensor `y` where all entries are sorted along the last dimension, in **ascending** order.

It also returns a Tensor `i` that provides the corresponding indices from `x`.

`y, i = torch.sort(x, d)` performs the sort operation along a specific dimension `d`.

`y, i = torch.sort(x)` is therefore equivalent to `y, i = torch.sort(x, x:dim())`

`y, i = torch.sort(x, d, true)` performs the sort operation along a specific dimension `d`, in **descending** order.

```

> x = torch.randn(3, 3)
> x

```

```

-1.2470 -0.4288 -0.5337
 0.8836 -0.1622  0.9604
 0.6297  0.2397  0.0746
[torch.DoubleTensor of size 3x3]

> torch.sort(x)
-1.2470 -0.5337 -0.4288
-0.1622  0.8836  0.9604
 0.0746  0.2397  0.6297
[torch.DoubleTensor of size 3x3]

 1  3  2
 2  1  3
 3  2  1
[torch.LongTensor of size 3x3]

```

`torch.topk([resval, resind,] x, k, [,dim] [,dir] [,sort])`

`y, i = torch.topk(x, k)` returns all `k` smallest elements in `x` over its last dimension including their indices, in unsorted order.

`y, i = torch.topk(x, k, dim)` performs the same operation except over dimension `dim`.

`y, i = torch.topk(x, k, dim, dir)` adds a sorting direction that has the same sense as `torch.sort`; `false` returns the `k` smallest elements in the slice, `true` returns the `k` largest elements in the slice.

`y, i = torch.topk(x, k, dim, dir, true)` specifies that the results in `y` should be sorted with respect to `dir`; by default, the results are potentially unsorted since the computation may be faster, but if sorting is desired, the sort flag may be passed, in which case the results are returned from smallest to `k`-th smallest (`dir == false`) or highest to `k`-th highest (`dir == true`).

The implementation provides no guarantee of the order of selection (indices) among equivalent elements (e.g., `torch.topk k == 2` selection of a vector `{1, 2, 1, 1}`; the values returned could be any pair of `1` entries in the vector).

`[res] torch.std([res,] x, [,dim] [,flag])`

`y = torch.std(x)` returns the standard deviation of the elements of `x`.

`y = torch.std(x, dim)` performs the `std` operation over the dimension `dim`.

`y = torch.std(x, dim, false)` performs the `std` operation normalizing by `n-1` (this is the default).

`y = torch.std(x, dim, true)` performs the `std` operation normalizing by `n` instead of `n-1`.

[res] torch.sum([res,] x)

`y = torch.sum(x)` returns the sum of the elements of `x`.

`y = torch.sum(x, 2)` performs the sum operation for each row.

`y = torch.sum(x, n)` performs the sum operation over the dimension `n`.

[res] torch.var([res,] x [,dim] [,flag])

`y = torch.var(x)` returns the variance of the elements of `x`.

`y = torch.var(x, dim)` performs the `var` operation over the dimension `dim`.

`y = torch.var(x, dim, false)` performs the `var` operation normalizing by `n-1` (this is the default).

`y = torch.var(x, dim, true)` performs the `var` operation normalizing by `n` instead of `n-1`.

Matrix-wide operations (Tensor-wide operations)

Note that many of the operations in [dimension-wise operations](#) can also be used as matrix-wide operations, by just omitting the `dim` parameter.

`torch.norm(x [,p] [,dim])`

`y = torch.norm(x)` returns the 2 -norm of the Tensor `x`.

`y = torch.norm(x, p)` returns the `p` -norm of the Tensor `x`.

`y = torch.norm(x, p, dim)` returns the `p` -norms of the Tensor `x` computed over the dimension `dim`.

`torch.renorm([res], x, p, dim, maxnorm)`

Renormalizes the sub- Tensor s along dimension `dim` such that they do not exceed norm `maxnorm`.

`y = torch.renorm(x, p, dim, maxnorm)` returns a version of `x` with `p` -norms lower than `maxnorm` over non- `dim` dimensions.

The `dim` argument is not to be confused with the argument of the same name in function [norm](#).

In this case, the `p` -norm is measured for each `i` -th sub- Tensor `x:select(dim, i)`. This function is equivalent to (but faster than) the following:

```
function renorm(matrix, value, dim, maxnorm)
  local m1 = matrix:transpose(dim, 1):contiguous()
  -- collapse non-dim dimensions:
  m2 = m1:reshape(m1:size(1), m1:nElement()/m1:size(1))
  local norms = m2:norm(value, 2)
  -- clip
  local new_norms = norms:clone()
  new_norms[torch.gt(norms, maxnorm)] = maxnorm
  new_norms:cdiv(norms:add(1e-7))
  -- renormalize
  m1:cmul(new_norms:expandAs(m1))
  return m1:transpose(dim, 1)
end
```

`x:renorm(p, dim, maxnorm)` returns the equivalent of `x:copy(torch.renorm(x, p, dim, maxnorm))`.

Note: this function is particularly useful as a regularizer for constraining the norm of parameter Tensor s.

See [Hinton et al. 2012, p. 2](#).

`torch.dist(x, y)`

`y = torch.dist(x, y)` returns the 2 -norm of $x - y$.

`y = torch.dist(x, y, p)` returns the p -norm of $x - y$.

`torch.numel(x)`

`y = torch.numel(x)` returns the count of the number of elements in the matrix x .

`torch.trace(x)`

`y = torch.trace(x)` returns the trace (sum of the diagonal elements) of a matrix x .

This is equal to the sum of the eigenvalues of x .

The returned value y is a number, not a `Tensor`.

Convolution Operations

These functions implement convolution or cross-correlation of an input image (or set of input images) with a kernel (or set of kernels).

The convolution function in Torch can handle different types of input/kernel dimensions and produces corresponding outputs.

The general form of operations always remain the same.

`[res] torch.conv2([res,] x, k, [, 'F' or 'V'])`

This function computes 2 dimensional convolutions between x and k .

These operations are similar to BLAS operations when number of dimensions of input and kernel are reduced by 2.

- x and k are 2D: convolution of a single image with a single kernel (2D output). This operation is similar to multiplication of two scalars.
- $x (p \times m \times n)$ and $k (p \times k_i \times k_j)$ are 3D: convolution of each input slice with corresponding kernel (3D output).
- $x (p \times m \times n)$ 3D, $k (q \times p \times k_i \times k_j)$ 4D: convolution of all input slices with the corresponding slice of kernel. Output is 3D $(q \times m \times n)$. This operation is similar to matrix vector product of matrix k and vector x .

The last argument controls if the convolution is a full ('F') or valid ('V') convolution. The default is **valid** convolution.

```
x = torch.rand(100, 100)
k = torch.rand(10, 10)
c = torch.conv2(x, k)
> c:size()
  91
  91
[torch.LongStorage of size 2]

c = torch.conv2(x, k, 'F')
> c:size()
 109
 109
[torch.LongStorage of size 2]
```

[res] torch.xcorr2([res,] x, k, [, 'F' or 'V'])

This function operates with same options and input/output configurations as `torch.conv2`, but performs cross-correlation of the input with the kernel k .

[res] torch.conv3([res,] x, k, [, 'F' or 'V'])

This function computes 3 dimensional convolutions between x and k .

These operations are similar to BLAS operations when number of dimensions of input and kernel are reduced by 3.

- x and k are 3D: convolution of a single image with a single kernel (3D output). This operation is similar to multiplication of two scalars.
- $x (p \times m \times n \times o)$ and $k (p \times k_i \times k_j \times k_k)$ are 4D: convolution of each input

slice with corresponding kernel (4D output).

- x ($p \times m \times n \times o$) 4D, k ($q \times p \times k_i \times k_j \times k_k$) 5D: convolution of all input slices with the corresponding slice of kernel. Output is 4D $q \times m \times n \times o$. This operation is similar to matrix vector product of matrix k and vector x .

The last argument controls if the convolution is a full ('F') or valid ('V') convolution. The default is **valid** convolution.

```
x = torch.rand(100, 100, 100)
k = torch.rand(10, 10, 10)
c = torch.conv3(x, k)
> c:size()
  91
  91
  91
[torch.LongStorage of size 3]

c = torch.conv3(x, k, 'F')
> c:size()
 109
 109
 109
[torch.LongStorage of size 3]
```

[res] torch.xcorr3([res,] x, k, [, 'F' or 'V'])

This function operates with same options and input/output configurations as `torch.conv3`, but performs cross-correlation of the input with the kernel k .

Eigenvalues, SVD, Linear System Solution

Functions in this section are implemented with an interface to [LAPACK](#) libraries.

If LAPACK libraries are not found during compilation step, then these functions will not be available.

[x, lu] torch.gesv([resb, resa,] B, A)

`X`, `LU = torch.gesv(B, A)` returns the solution of $AX = B$ and `LU` contains `L` and `U` factors for LU factorization of `A`.

`A` has to be a square and non-singular matrix (2D Tensor).

`A` and `LU` are $m \times m$, `X` is $m \times k$ and `B` is $m \times k$.

If `resb` and `resa` are given, then they will be used for temporary storage and returning the result.

- `resa` will contain `L` and `U` factors for LU factorization of `A`.
- `resb` will contain the solution `X`.

Note: Irrespective of the original strides, the returned matrices `resb` and `resa` will be transposed, i.e. with strides `1, m` instead of `m, 1`.

```
> a = torch.Tensor([
    [6.80, -2.11, 5.66, 5.97, 8.23],
    [-6.05, -3.30, 5.36, -4.44, 1.08],
    [-0.45, 2.58, -2.70, 0.27, 9.04],
    [8.32, 2.71, 4.35, -7.17, 2.14],
    [-9.67, -5.14, -7.26, 6.08, -6.87]]) :t()

> b = torch.Tensor([
    [4.02, 6.19, -8.22, -7.57, -3.03],
    [-1.56, 4.00, -8.67, 1.75, 2.86],
    [9.81, -4.09, -4.57, -8.61, 8.99]]) :t()

> b
 4.0200 -1.5600  9.8100
 6.1900  4.0000 -4.0900
-8.2200 -8.6700 -4.5700
-7.5700  1.7500 -8.6100
-3.0300  2.8600  8.9900
[torch.DoubleTensor of dimension 5x3]

> a
 6.8000 -6.0500 -0.4500  8.3200 -9.6700
-2.1100 -3.3000  2.5800  2.7100 -5.1400
 5.6600  5.3600 -2.7000  4.3500 -7.2600
 5.9700 -4.4400  0.2700 -7.1700  6.0800
 8.2300  1.0800  9.0400  2.1400 -6.8700
[torch.DoubleTensor of dimension 5x5]

> x = torch.gesv(b, a)
> x
-0.8007 -0.3896  0.9555
```



```

-0.6952 -0.5544  0.2207
 0.5939  0.8422  1.9006
 1.3217 -0.1038  5.3577
 0.5658  0.1057  4.0406
[torch.DoubleTensor of dimension 5x3]

> b:dist(a * x)
1.1682163181673e-14

```

[x] torch.trtrs([resb, resa,] b, a [, 'U' or 'L'] [, 'N' or 'T'] [, 'N' or 'U'])

$X = \text{torch.trtrs}(B, A)$ returns the solution of $AX = B$ where A is upper-triangular.

A has to be a square, triangular, non-singular matrix (2D Tensor).

A and $resa$ are $m \times m$, X and B are $m \times k$.

(To be very precise: A does not have to be triangular and non-singular, rather only its upper or lower triangle will be taken into account and that part has to be non-singular.)

The function has several options:

- `uplo` ('U' or 'L') specifies whether A is upper or lower triangular; the default value is 'U'.
- `trans` ('N' or 'T') specifies the system of equations: 'N' for $A * X = B$ (no transpose), or 'T' for $A^T * X = B$ (transpose); the default value is 'N'.
- `diag` ('N' or 'U') 'U' specifies that A is unit triangular, i.e., it has ones on its diagonal; 'N' specifies that A is not (necessarily) unit triangular; the default value is 'N'.

If $resb$ and $resa$ are given, then they will be used for temporary storage and returning the result.

$resb$ will contain the solution X .

Note: Irrespective of the original strides, the returned matrices $resb$ and $resa$ will be transposed, i.e. with strides 1, m instead of m , 1.

```

> a = torch.Tensor({{6.80, -2.11, 5.66, 5.97, 8.23},
                    {0, -3.30, 5.36, -4.44, 1.08},
                    {0, 0, -2.70, 0.27, 9.04},
                    {0, 0, 0, -7.17, 2.14},
                    {0, 0, 0, 0, -6.87}})

```

```

> b = torch.Tensor([4.02, 6.19, -8.22, -7.57, -3.03],
                    [-1.56, 4.00, -8.67, 1.75, 2.86],
                    [9.81, -4.09, -4.57, -8.61, 8.99]):t()

> b
 4.0200 -1.5600  9.8100
 6.1900  4.0000 -4.0900
-8.2200 -8.6700 -4.5700
-7.5700  1.7500 -8.6100
-3.0300  2.8600  8.9900
[torch.DoubleTensor of dimension 5x3]

> a
 6.8000 -2.1100  5.6600  5.9700  8.2300
 0.0000 -3.3000  5.3600 -4.4400  1.0800
 0.0000  0.0000 -2.7000  0.2700  9.0400
 0.0000  0.0000  0.0000 -7.1700  2.1400
 0.0000  0.0000  0.0000  0.0000 -6.8700
[torch.DoubleTensor of dimension 5x5]

> x = torch.trtrs(b, a)
> x
-3.5416 -0.2514  3.0847
 4.2072  2.0391 -4.5146
 4.6399  1.7804 -2.6077
 1.1874 -0.3683  0.8103
 0.4410 -0.4163 -1.3086
[torch.DoubleTensor of size 5x3]

> b:dist(a*x)
4.1895292266754e-15

```

torch.potrf([res,] A [, 'U' or 'L'])

Cholesky Decomposition of 2D Tensor A.

The matrix A has to be a positive-definite and either symmetric or complex Hermitian.

The factorization has the form

$$\begin{aligned}
 A &= U^*T * U, & \text{if } \text{UPLO} = 'U', & \text{ or} \\
 A &= L * L^*T, & \text{if } \text{UPLO} = 'L', &
 \end{aligned}$$

where `U` is an upper triangular matrix and `L` is lower triangular.

The optional character `uplo` = {'U', 'L'} specifies whether the upper or lower triangular decomposition should be returned. By default, `uplo` = 'U'.

`U = torch.potrf(A, 'U')` returns the upper triangular Cholesky decomposition of `A`.

`L = torch.potrf(A, 'L')` returns the lower triangular Cholesky decomposition of `A`.

If Tensor `res` is provided, the resulting decomposition will be stored therein.

```
> A = torch.Tensor({
  {1.2705,  0.9971,  0.4948,  0.1389,  0.2381},
  {0.9971,  0.9966,  0.6752,  0.0686,  0.1196},
  {0.4948,  0.6752,  1.1434,  0.0314,  0.0582},
  {0.1389,  0.0686,  0.0314,  0.0270,  0.0526},
  {0.2381,  0.1196,  0.0582,  0.0526,  0.3957}})

> chol = torch.potrf(A)
> chol
 1.1272  0.8846  0.4390  0.1232  0.2112
 0.0000  0.4626  0.6200 -0.0874 -0.1453
 0.0000  0.0000  0.7525  0.0419  0.0738
 0.0000  0.0000  0.0000  0.0491  0.2199
 0.0000  0.0000  0.0000  0.0000  0.5255
[torch.DoubleTensor of size 5x5]

> torch.potrf(chol, A, 'L')
> chol
 1.1272  0.0000  0.0000  0.0000  0.0000
 0.8846  0.4626  0.0000  0.0000  0.0000
 0.4390  0.6200  0.7525  0.0000  0.0000
 0.1232 -0.0874  0.0419  0.0491  0.0000
 0.2112 -0.1453  0.0738  0.2199  0.5255
[torch.DoubleTensor of size 5x5]
```

`torch.pstrf([res, piv,] A [, 'U' or 'L'])`

Cholesky factorization with complete pivoting of a real symmetric positive semidefinite 2D Tensor `A`.

The matrix `A` has to be a positive semi-definite and symmetric. The factorization has the form

$$P^{**T} * A * P = U^{**T} * U, \quad \text{if } \text{UPLO} = 'U',$$

$$P^{**T} * A * P = L * L^{**T}, \quad \text{if } \text{UPLO} = 'L',$$

where U is an upper triangular matrix and L is lower triangular, and P is stored as the vector `piv`. More specifically, `piv` is such that the nonzero entries are $P[\text{piv}[k], k] = 1$.

The optional character argument `uplo` = {'U', 'L'} specifies whether the upper or lower triangular decomposition should be returned. By default, `uplo` = 'U'.

`U, piv = torch.sdtrf(A, 'U')` returns the upper triangular Cholesky decomposition of A

`L, piv = torch.potrf(A, 'L')` returns the lower triangular Cholesky decomposition of A .

If tensors `res` and `piv` (an `IntTensor`) are provided, the resulting decomposition will be stored therein.

```
> A = torch.Tensor({
    {1.2705, 0.9971, 0.4948, 0.1389, 0.2381},
    {0.9971, 0.9966, 0.6752, 0.0686, 0.1196},
    {0.4948, 0.6752, 1.1434, 0.0314, 0.0582},
    {0.1389, 0.0686, 0.0314, 0.0270, 0.0526},
    {0.2381, 0.1196, 0.0582, 0.0526, 0.3957}}})
```

```
> U, piv = torch.pstrf(A)
> U
1.1272 0.4390 0.2112 0.8846 0.1232
0.0000 0.9750 -0.0354 0.2942 -0.0233
0.0000 0.0000 0.5915 -0.0961 0.0435
0.0000 0.0000 0.0000 0.3439 -0.0854
0.0000 0.0000 0.0000 0.0000 0.0456
[torch.DoubleTensor of size 5x5]
```

```
> piv
1
3
5
2
4
[torch.IntTensor of size 5]
```

```
> Ap = U:t() * U
```

```

> Ap
1.2705  0.4948  0.2381  0.9971  0.1389
0.4948  1.1434  0.0582  0.6752  0.0314
0.2381  0.0582  0.3957  0.1196  0.0526
0.9971  0.6752  0.1196  0.9966  0.0686
0.1389  0.0314  0.0526  0.0686  0.0270
[torch.DoubleTensor of size 5x5]

> -- Permute rows and columns
> Ap:indexCopy(1, piv:long(), Ap:clone())
> Ap:indexCopy(2, piv:long(), Ap:clone())
> (Ap - A):norm()
1.5731560566382e-16

```

torch.potrs([res,] B, chol [, 'U' or 'L'])

Returns the solution to linear system $AX = B$ using the Cholesky decomposition `chol` of 2D Tensor `A`.

Square matrix `chol` should be triangular; and, righthand side matrix `B` should be of full rank.

Optional character `uplo` = {'U', 'L'} specifies matrix `chol` as either upper or lower triangular; and, by default, equals 'U'.

If Tensor `res` is provided, the resulting decomposition will be stored therein.

```

> A = torch.Tensor({
    {1.2705,  0.9971,  0.4948,  0.1389,  0.2381},
    {0.9971,  0.9966,  0.6752,  0.0686,  0.1196},
    {0.4948,  0.6752,  1.1434,  0.0314,  0.0582},
    {0.1389,  0.0686,  0.0314,  0.0270,  0.0526},
    {0.2381,  0.1196,  0.0582,  0.0526,  0.3957}})

> B = torch.Tensor({
    {0.6219,  0.3439,  0.0431},
    {0.5642,  0.1756,  0.0153},
    {0.2334,  0.8594,  0.4103},
    {0.7556,  0.1966,  0.9637},
    {0.1420,  0.7185,  0.7476}})

> chol = torch.potrf(A)
> chol

```

```

1.1272  0.8846  0.4390  0.1232  0.2112
0.0000  0.4626  0.6200 -0.0874 -0.1453
0.0000  0.0000  0.7525  0.0419  0.0738
0.0000  0.0000  0.0000  0.0491  0.2199
0.0000  0.0000  0.0000  0.0000  0.5255

```

```
[torch.DoubleTensor of size 5x5]
```

```
> solve = torch.potrs(B, chol)
```

```
> solve
```

```

12.1945   61.8622   92.6882
-11.1782  -97.0303 -138.4874
-15.3442  -76.6562 -116.8218
   6.1930   13.5238   25.2056
  29.9678  251.7346  360.2301

```

```
[torch.DoubleTensor of size 5x3]
```

```
> A*solve
```

```

0.6219  0.3439  0.0431
0.5642  0.1756  0.0153
0.2334  0.8594  0.4103
0.7556  0.1966  0.9637
0.1420  0.7185  0.7476

```

```
[torch.DoubleTensor of size 5x3]
```

```
> B:dist(A*solve)
```

```
4.6783066076306e-14
```

torch.potri([res,] chol [, 'U' or 'L'])

Returns the inverse of 2D `Tensor` `A` given its Cholesky decomposition `chol`.

Square matrix `chol` should be triangular.

Optional character `uplo` = {'U', 'L'} specifies matrix `chol` as either upper or lower triangular; and, by default, equals 'U'.

If `Tensor` `res` is provided, the resulting inverse will be stored therein.

```

> A = torch.Tensor({
    {1.2705,  0.9971,  0.4948,  0.1389,  0.2381},
    {0.9971,  0.9966,  0.6752,  0.0686,  0.1196},
    {0.4948,  0.6752,  1.1434,  0.0314,  0.0582},

```

```

    {0.1389, 0.0686, 0.0314, 0.0270, 0.0526},
    {0.2381, 0.1196, 0.0582, 0.0526, 0.3957}})

> chol = torch.potrf(A)
> chol
1.1272  0.8846  0.4390  0.1232  0.2112
0.0000  0.4626  0.6200 -0.0874 -0.1453
0.0000  0.0000  0.7525  0.0419  0.0738
0.0000  0.0000  0.0000  0.0491  0.2199
0.0000  0.0000  0.0000  0.0000  0.5255
[torch.DoubleTensor of size 5x5]

> inv = torch.potri(chol)
> inv
42.2781  -39.0824   8.3019 -133.4998   2.8980
-39.0824   38.1222  -8.7468  119.4247  -2.5944
 8.3019   -8.7468   3.1104 -25.1405   0.5327
-133.4998  119.4247 -25.1405  480.7511 -15.9747
 2.8980   -2.5944   0.5327 -15.9747   3.6127
[torch.DoubleTensor of size 5x5]

> inv:dist(torch.inverse(A))
2.8525852877633e-12

```

torch.gels([resb, resa,] b, a)

Solution of least squares and least norm problems for a full rank $m \times n$ matrix A .

- If $n \leq m$, then solve $\|AX - B\|_F$.
- If $n > m$, then solve $\min \|X\|_F$ s.t. $AX = B$.

On return, first n rows of x matrix contains the solution and the rest contains residual information.

Square root of sum squares of elements of each column of x starting at row $n + 1$ is the residual for corresponding column.

Note: Irrespective of the original strides, the returned matrices `resb` and `resa` will be transposed, i.e. with strides `1, m` instead of `m, 1`.

```

> a = torch.Tensor([
    [ 1.44, -9.96, -7.55,  8.34,  7.08, -5.45],
    [-7.84, -0.28,  3.24,  8.09,  2.52, -5.70],
    [-4.39, -3.24,  6.27,  5.28,  0.74, -1.19],

```

```

        {4.53,  3.83, -6.64,  2.06, -2.47,  4.70}}):t()

> b = torch.Tensor({{8.58,  8.26,  8.48, -5.28,  5.72,  8.93},
                    {9.35, -4.43, -0.70, -0.26, -7.36, -2.52}}):t()

> a
  1.4400 -7.8400 -4.3900  4.5300
 -9.9600 -0.2800 -3.2400  3.8300
 -7.5500  3.2400  6.2700 -6.6400
  8.3400  8.0900  5.2800  2.0600
  7.0800  2.5200  0.7400 -2.4700
 -5.4500 -5.7000 -1.1900  4.7000
[torch.DoubleTensor of dimension 6x4]

> b
  8.5800  9.3500
  8.2600 -4.4300
  8.4800 -0.7000
 -5.2800 -0.2600
  5.7200 -7.3600
  8.9300 -2.5200
[torch.DoubleTensor of dimension 6x2]

> x = torch.gels(b, a)
> x
 -0.4506  0.2497
 -0.8492 -0.9020
  0.7066  0.6323
  0.1289  0.1351
 13.1193 -7.4922
 -4.8214 -7.1361
[torch.DoubleTensor of dimension 6x2]

> b:dist(a*x:narrow(1, 1, 4))
17.390200628863

> math.sqrt(x:narrow(1, 5, 2):pow(2):sumall())
17.390200628863

```

`torch.symeig([rese, resv,] a [, 'N' or 'V'] [, 'U' or 'L'])`

`e, V = torch.symeig(A)` returns eigenvalues and eigenvectors of a symmetric real matrix

A.

A and V are $m \times m$ matrices and e is a m dimensional vector.

This function calculates all eigenvalues (and vectors) of A such that $A = V \text{diag}(e) V'$.

Third argument defines computation of eigenvectors or eigenvalues only.

If it is 'N', only eigenvalues are computed.

If it is 'V', both eigenvalues and eigenvectors are computed.

Since the input matrix A is supposed to be symmetric, only upper triangular portion is used by default.

If the 4th argument is 'L', then lower triangular portion is used.

Note: Irrespective of the original strides, the returned matrix V will be transposed, i.e. with strides 1, m instead of m, 1.

```
> a = torch.Tensor({{ 1.96, 0.00, 0.00, 0.00, 0.00},
                    {-6.49, 3.80, 0.00, 0.00, 0.00},
                    {-0.47, -6.39, 4.17, 0.00, 0.00},
                    {-7.20, 1.50, -1.51, 5.70, 0.00},
                    {-0.65, -6.34, 2.67, 1.80, -7.10}}):t()
```

```
> a
 1.9600 -6.4900 -0.4700 -7.2000 -0.6500
 0.0000  3.8000 -6.3900  1.5000 -6.3400
 0.0000  0.0000  4.1700 -1.5100  2.6700
 0.0000  0.0000  0.0000  5.7000  1.8000
 0.0000  0.0000  0.0000  0.0000 -7.1000
[torch.DoubleTensor of dimension 5x5]
```

```
> e = torch.symeig(a)
> e
-11.0656
 -6.2287
  0.8640
  8.8655
 16.0948
[torch.DoubleTensor of dimension 5]
```

```
> e, v = torch.symeig(a, 'V')
> e
-11.0656
 -6.2287
  0.8640
```

```

8.8655
16.0948
[torch.DoubleTensor of dimension 5]

> v
-0.2981 -0.6075  0.4026 -0.3745  0.4896
-0.5078 -0.2880 -0.4066 -0.3572 -0.6053
-0.0816 -0.3843 -0.6600  0.5008  0.3991
-0.0036 -0.4467  0.4553  0.6204 -0.4564
-0.8041  0.4480  0.1725  0.3108  0.1622
[torch.DoubleTensor of dimension 5x5]

> v*torch.diag(e)*v:t()
1.9600 -6.4900 -0.4700 -7.2000 -0.6500
-6.4900  3.8000 -6.3900  1.5000 -6.3400
-0.4700 -6.3900  4.1700 -1.5100  2.6700
-7.2000  1.5000 -1.5100  5.7000  1.8000
-0.6500 -6.3400  2.6700  1.8000 -7.1000
[torch.DoubleTensor of dimension 5x5]

> a:dist(torch.triu(v*torch.diag(e)*v:t()))
1.0219480822443e-14

```

torch.eig([rese, resv,] a [, 'N' or 'V'])

`e, V = torch.eig(A)` returns eigenvalues and eigenvectors of a general real square matrix `A`.

`A` and `V` are $m \times m$ matrices and `e` is a m dimensional vector.

This function calculates all right eigenvalues (and vectors) of `A` such that $A = V \text{diag}(e) V^T$.

Third argument defines computation of eigenvectors or eigenvalues only.

If it is `'N'`, only eigenvalues are computed.

If it is `'V'`, both eigenvalues and eigenvectors are computed.

The eigen values returned follow [LAPACK convention](#) and are returned as complex (real/imaginary) pairs of numbers ($2 * m$ dimensional Tensor).

Note: Irrespective of the original strides, the returned matrix `V` will be transposed, i.e. with strides `1, m` instead of `m, 1`.

```
> a = torch.Tensor([[ 1.96,  0.00,  0.00,  0.00,  0.00],
                    [-6.49,  3.80,  0.00,  0.00,  0.00],
                    [-0.47, -6.39,  4.17,  0.00,  0.00],
                    [-7.20,  1.50, -1.51,  5.70,  0.00],
                    [-0.65, -6.34,  2.67,  1.80, -7.10]]):t()
```

```
> a
1.9600 -6.4900 -0.4700 -7.2000 -0.6500
0.0000  3.8000 -6.3900  1.5000 -6.3400
0.0000  0.0000  4.1700 -1.5100  2.6700
0.0000  0.0000  0.0000  5.7000  1.8000
0.0000  0.0000  0.0000  0.0000 -7.1000
[torch.DoubleTensor of dimension 5x5]
```

```
> b = a + torch.triu(a, 1):t()
> b
```

```
1.9600 -6.4900 -0.4700 -7.2000 -0.6500
-6.4900  3.8000 -6.3900  1.5000 -6.3400
-0.4700 -6.3900  4.1700 -1.5100  2.6700
-7.2000  1.5000 -1.5100  5.7000  1.8000
-0.6500 -6.3400  2.6700  1.8000 -7.1000
[torch.DoubleTensor of dimension 5x5]
```

```
> e = torch.eig(b)
> e
16.0948  0.0000
-11.0656  0.0000
-6.2287  0.0000
 0.8640  0.0000
 8.8655  0.0000
[torch.DoubleTensor of dimension 5x2]
```

```
> e, v = torch.eig(b, 'V')
> e
16.0948  0.0000
-11.0656  0.0000
-6.2287  0.0000
 0.8640  0.0000
 8.8655  0.0000
[torch.DoubleTensor of dimension 5x2]
```

```
> v
-0.4896  0.2981 -0.6075 -0.4026 -0.3745
 0.6053  0.5078 -0.2880  0.4066 -0.3572
```

```

-0.3991  0.0816 -0.3843  0.6600  0.5008
 0.4564  0.0036 -0.4467 -0.4553  0.6204
-0.1622  0.8041  0.4480 -0.1725  0.3108
[torch.DoubleTensor of dimension 5x5]

> v * torch.diag(e:select(2, 1))*v:t()
 1.9600 -6.4900 -0.4700 -7.2000 -0.6500
-6.4900  3.8000 -6.3900  1.5000 -6.3400
-0.4700 -6.3900  4.1700 -1.5100  2.6700
-7.2000  1.5000 -1.5100  5.7000  1.8000
-0.6500 -6.3400  2.6700  1.8000 -7.1000
[torch.DoubleTensor of dimension 5x5]

> b:dist(v * torch.diag(e:select(2, 1)) * v:t())
3.5423944346685e-14

```

torch.svd([resu, ress, resv,] a [, 'S' or 'A'])

`U, S, V = torch.svd(A)` returns the singular value decomposition of a real matrix `A` of size `n × m` such that $A = USV^*$.

`U` is `n × n`, `S` is `n × m` and `V` is `m × m`.

The last argument, if it is string, represents the number of singular values to be computed.

'S' stands for *some* and 'A' stands for *all*.

Note: Irrespective of the original strides, the returned matrix `U` will be transposed, i.e. with strides `1, n` instead of `n, 1`.

```

> a = torch.Tensor({{8.79,  6.11, -9.15,  9.57, -3.49,  9.84},
                    {9.93,  6.91, -7.93,  1.64,  4.02,  0.15},
                    {9.83,  5.04,  4.86,  8.83,  9.80, -8.99},
                    {5.45, -0.27,  4.85,  0.74, 10.00, -6.02},
                    {3.16,  7.98,  3.01,  5.80,  4.27, -5.31}}):t()

> a
 8.7900  9.9300  9.8300  5.4500  3.1600
 6.1100  6.9100  5.0400 -0.2700  7.9800
-9.1500 -7.9300  4.8600  4.8500  3.0100
 9.5700  1.6400  8.8300  0.7400  5.8000
-3.4900  4.0200  9.8000 10.0000  4.2700
 9.8400  0.1500 -8.9900 -6.0200 -5.3100

```

```

> u, s, v = torch.svd(a)
> u
-0.5911  0.2632  0.3554  0.3143  0.2299
-0.3976  0.2438 -0.2224 -0.7535 -0.3636
-0.0335 -0.6003 -0.4508  0.2334 -0.3055
-0.4297  0.2362 -0.6859  0.3319  0.1649
-0.4697 -0.3509  0.3874  0.1587 -0.5183
 0.2934  0.5763 -0.0209  0.3791 -0.6526
[torch.DoubleTensor of dimension 6x5]

> s
27.4687
22.6432
 8.5584
 5.9857
 2.0149
[torch.DoubleTensor of dimension 5]

> v
-0.2514  0.8148 -0.2606  0.3967 -0.2180
-0.3968  0.3587  0.7008 -0.4507  0.1402
-0.6922 -0.2489 -0.2208  0.2513  0.5891
-0.3662 -0.3686  0.3859  0.4342 -0.6265
-0.4076 -0.0980 -0.4933 -0.6227 -0.4396
[torch.DoubleTensor of dimension 5x5]

> u * torch.diag(s) * v:t()
 8.7900  9.9300  9.8300  5.4500  3.1600
 6.1100  6.9100  5.0400 -0.2700  7.9800
-9.1500 -7.9300  4.8600  4.8500  3.0100
 9.5700  1.6400  8.8300  0.7400  5.8000
-3.4900  4.0200  9.8000 10.0000  4.2700
 9.8400  0.1500 -8.9900 -6.0200 -5.3100
[torch.DoubleTensor of dimension 6x5]

> a:dist(u * torch.diag(s) * v:t())
2.8923773593204e-14

```

torch.inverse([res,] x)

Computes the inverse of square matrix `x`.

`torch.inverse(x)` returns the result as a new matrix.

`torch.inverse(y, x)` puts the result in `y`.

Note: Irrespective of the original strides, the returned matrix `y` will be transposed, i.e. with strides `1, m` instead of `m, 1`.

```
> x = torch.rand(10, 10)
> y = torch.inverse(x)
> z = x * y
> z
 1.0000 -0.0000  0.0000 -0.0000  0.0000  0.0000  0.0000 -0.0000
 0.0000  0.0000
 0.0000  1.0000 -0.0000 -0.0000  0.0000  0.0000 -0.0000 -0.0000
-0.0000  0.0000
 0.0000 -0.0000  1.0000 -0.0000  0.0000  0.0000 -0.0000 -0.0000
 0.0000  0.0000
 0.0000 -0.0000 -0.0000  1.0000 -0.0000  0.0000  0.0000 -0.0000
-0.0000  0.0000
 0.0000 -0.0000  0.0000 -0.0000  1.0000  0.0000  0.0000 -0.0000
-0.0000  0.0000
 0.0000 -0.0000  0.0000 -0.0000  0.0000  1.0000  0.0000 -0.0000
-0.0000  0.0000
 0.0000 -0.0000  0.0000 -0.0000  0.0000  0.0000  1.0000 -0.0000
 0.0000  0.0000
 0.0000 -0.0000 -0.0000 -0.0000  0.0000  0.0000  0.0000  1.0000
 0.0000  0.0000
 0.0000 -0.0000 -0.0000 -0.0000  0.0000  0.0000 -0.0000 -0.0000
 1.0000  0.0000
 0.0000 -0.0000  0.0000 -0.0000  0.0000  0.0000  0.0000 -0.0000
 0.0000  1.0000
[torch.DoubleTensor of dimension 10x10]

> torch.max(torch.abs(z - torch.eye(10))) -- Max nonzero
2.3092638912203e-14
```

`torch.qr([q, r], x)`

Compute a QR decomposition of the matrix `x`: matrices `q` and `r` such that `x = q * r`, with `q` orthogonal and `r` upper triangular.

This returns the thin (reduced) QR factorization.

`torch.qr(x)` returns the Q and R components as new matrices.

`torch.qr(q, r, x)` stores them in existing Tensor s `q` and `r`.

Note that precision may be lost if the magnitudes of the elements of `x` are large.

Note also that, while it should always give you a valid decomposition, it may not give you the same one across platforms - it will depend on your LAPACK implementation.

Note: Irrespective of the original strides, the returned matrix `q` will be transposed, i.e. with strides `1, m` instead of `m, 1`.

```
> a = torch.Tensor({12, -51, 4}, {6, 167, -68}, {-4, 24, -41})
> a
  12  -51   4
   6  167 -68
  -4   24 -41
[torch.DoubleTensor of dimension 3x3]

> q, r = torch.qr(a)
> q
-0.8571  0.3943  0.3314
-0.4286 -0.9029 -0.0343
 0.2857 -0.1714  0.9429
[torch.DoubleTensor of dimension 3x3]

> r
-14.0000 -21.0000  14.0000
  0.0000 -175.0000  70.0000
  0.0000   0.0000 -35.0000
[torch.DoubleTensor of dimension 3x3]

> (q * r):round()
  12  -51   4
   6  167 -68
  -4   24 -41
[torch.DoubleTensor of dimension 3x3]

> (q:t() * q):round()
  1  0  0
  0  1  0
  0  0  1
[torch.DoubleTensor of dimension 3x3]
```

`torch.geqrf([m, tau], a)`

This is a low-level function for calling LAPACK directly.
You'll generally want to use `torch.qr()` instead.

Computes a QR decomposition of `a`, but without constructing Q and R as explicit separate matrices.

Rather, this directly calls the underlying LAPACK function `?geqrf` which produces a sequence of 'elementary reflectors'.

See [LAPACK documentation](#) for further details.

`torch.orgqr([q], m, tau)`

This is a low-level function for calling LAPACK directly.
You'll generally want to use `torch.qr()` instead.

Constructs a Q matrix from a sequence of elementary reflectors, such as that given by `torch.geqrf`.

See [LAPACK documentation](#) for further details.

`torch.ormqr([res], m, tau, mat [, 'L' or 'R'] [, 'N' or 'T'])`

Multiply a matrix with `Q` as defined by the elementary reflectors and scalar factors returned by `geqrf`.

This is a low-level function for calling LAPACK directly.
You'll generally want to use `torch.qr()` instead.

- `side` ('L' or 'R') specifies whether `mat` should be left-multiplied, `mat * Q`, or right-multiplied, `Q * mat`.
- `trans` ('N' or 'T') specifies whether `Q` should be transposed before being multiplied.

See [LAPACK documentation](#) for further details.

Logical Operations on `Tensors`

These functions implement logical comparison operators that take a `Tensor` as input and another `Tensor` or a number as the comparison target. They return a `ByteTensor` in which each element is `0` or `1` indicating if the comparison for the corresponding element was `false` or `true` respectively.

`torch.lt(a, b)`

Implements `<` operator comparing each element in `a` with `b` (if `b` is a number) or each element in `a` with corresponding element in `b`.

`torch.le(a, b)`

Implements `<=` operator comparing each element in `a` with `b` (if `b` is a number) or each element in `a` with corresponding element in `b`.

`torch.gt(a, b)`

Implements `>` operator comparing each element in `a` with `b` (if `b` is a number) or each element in `a` with corresponding element in `b`.

`torch.ge(a, b)`

Implements `>=` operator comparing each element in `a` with `b` (if `b` is a number) or each element in `a` with corresponding element in `b`.

`torch.eq(a, b)`

Implements `==` operator comparing each element in `a` with `b` (if `b` is a number) or each element in `a` with corresponding element in `b`.

`torch.ne(a, b)`

Implements `~=` operator comparing each element in `a` with `b` (if `b` is a number) or each element in `a` with corresponding element in `b`.

`torch.all(a)`

`torch.any(a)`

Additionally, `any` and `all` logically sum a `ByteTensor` returning `true` if any or all elements are logically true respectively.

Note that logically true here is meant in the C sense (zero is `false`, non-zero is `true`) such as the output of the `Tensor` element-wise logical operations.

```
> a = torch.rand(10)
> b = torch.rand(10)
> a
0.5694
0.5264
0.3041
0.4159
0.1677
0.7964
0.0257
0.2093
0.6564
0.0740
[torch.DoubleTensor of dimension 10]

> b
0.2950
0.4867
0.9133
0.1291
0.1811
0.3921
0.7750
0.3259
0.2263
0.1737
[torch.DoubleTensor of dimension 10]
```

```
> torch.lt(a, b)
0
0
1
0
1
0
1
1
0
1
[torch.ByteTensor of dimension 10]
```

```
> torch.eq(a, b)
0
0
0
0
0
0
0
0
0
0
[torch.ByteTensor of dimension 10]
```

```
> torch.ne(a, b)
1
1
1
1
1
1
1
1
1
1
[torch.ByteTensor of dimension 10]
```

```
> torch.gt(a, b)
1
1
0
1
```

```
0
1
0
0
1
0
[torch.ByteTensor of dimension 10]

> a[torch.gt(a, b)] = 10
> a
10.0000
10.0000
 0.3041
10.0000
 0.1677
10.0000
 0.0257
 0.2093
10.0000
 0.0740
[torch.DoubleTensor of dimension 10]

> a[torch.gt(a, 1)] = -1
> a
-1.0000
-1.0000
 0.3041
-1.0000
 0.1677
-1.0000
 0.0257
 0.2093
-1.0000
 0.0740
[torch.DoubleTensor of dimension 10]

> a = torch.ones(3):byte()
> torch.all(a)
true

> a[2] = 0
> torch.all(a)
false

> torch.any(a)
```

true

```
> a:zero()
```

```
> torch.any(a)
```

false